

GAN Training and Inference

December 11, 2021

```
[ ]: # Updated GAN CODE
# Data pre-processing pipeline + GAN for the KITTI Dataset.
import torch
from torchvision import transforms
import numpy as np
import cv2
from glob import glob
import os
import matplotlib.pyplot as plt
from skimage import io
import torchvision.models as models
import torch
import numpy as np
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
from torch.autograd import Variable
import cv2
import matplotlib.pyplot as plt
from tqdm import tqdm
from glob import glob
import numpy as np
import torch.nn.init
import os
import torch
import numpy as np
```

```
[ ]: from google.colab import drive
drive.mount(r'/content/drive/', force_remount = True) # mounting the drive.
```

```
[ ]: basePath = r'/media/athrva/New Volume/cis520/Final Project/Processed_Data'
```

```
[ ]: # Some Custom Transformations (adapted code)

import torch
import random
```

```

import numpy as np
from PIL import Image

'''Set of transform random routines that takes list of inputs as arguments,
in order to have random but coherent transformations.'''

class Compose(object):
    def __init__(self, transforms):
        self.transforms = transforms

    def __call__(self, images):
        for t in self.transforms:
            images = t(images)
        return images

class Normalize(object):
    def __init__(self, mean, std):
        self.mean = mean
        self.std = std

    def __call__(self, images):
        for tensor in images:
            for t, m, s in zip(tensor, self.mean, self.std):
                t.sub_(m).div_(s)
        return images

class ArrayToTensor(object):
    def __call__(self, images):
        tensors = []
        for im in images:
            # put it from HWC to CHW format
            im = np.transpose(im, (2, 0, 1))
            tensors.append(torch.from_numpy(im).float()/255)
        return tensors

```

```

[ ]: ### Dataloaders using pytorch frameworks
from torch.utils.data import Dataset, DataLoader
from skimage.transform import resize

def depthByName(path):
    num = path.split('.')[0].split(os.path.sep)[-1]
    return num

```

```

def imageByName(path):
    num = path.split('.jpg')[0].split(os.path.sep)[-1]
    return num

def sortPaths(imgPaths, dmapPaths):
    imgPaths = sorted(imgPaths, key = imageByName);
    dmapPaths = sorted(dmapPaths, key = depthByName);
    return imgPaths, dmapPaths

def validityCheckPairs(img_dmap_pairs): # checks the validity of the incoming
    ↪data (checks whether data is pairwise)
    counter = 0

    for pair in img_dmap_pairs:
        img_name = pair[0].split('/')[0].split('.')[0].split('_')
        dmap_name = pair[1].split('/')[0].split('.')[0].split('_')

        img_ID = img_name[1] + img_name[3]
        dmap_ID = dmap_name[1] + dmap_name[3]

        assert(img_ID == dmap_ID)

        counter += 1

    print(f'Test Passed : All data is pairwise. Number of Samples tested : {
    ↪counter}')
    print('=====')

def buildPairs(dmaps, imgs):
    pairs = []
    for i in imgs:
        img_name = i.split('/')[0].split('.')[0].split('_')
        dmap_path = basePath + '/dmaps/' + img_name[0] + '_' + img_name[1] + \
            '_' + 'dmap' + '_' + img_name[3] + '.jpg'

        if dmap_path in dmaps:
            pairs.append([i, dmap_path])
    return pairs

def getDataPaths(): # Gets the paths where data is stored.
    Directories = glob(basePath + os.path.sep + '*')

    imgPathsDir = glob(Directories[0] + os.path.sep + '*.jpg');
    dmapPathsDir = glob(Directories[1] + os.path.sep + '*.jpg');
    imgPathsDir, dmapPathsDir = sortPaths(imgPathsDir, dmapPathsDir);

```

```

# Figure out which of the two has fewer entries
# bottleneck = np.argmin([len(imgPathsDir), len(dmapPathsDir)])
# For now, we know that images has fewer entries than dmaps
img_dmap_pairs = buildPairs(imgPathsDir, dmapPathsDir)

validityCheckPairs(img_dmap_pairs)
# Check again to be sure
validityCheckPairs(img_dmap_pairs)

return img_dmap_pairs

def rgb2gray(rgb):

    r, g, b = rgb[:, :, 0], rgb[:, :, 1], rgb[:, :, 2]
    gray = 0.2989 * r + 0.5870 * g + 0.1140 * b

    return gray

## Image Dataloader
class ImageDataset(Dataset):

    def __init__(self,
                  paths,
                  op,
                  transforms=None):
        """
        Args:
            op (str): "train", "val", or "test" to indicate the split type
            transforms (list or None): Image transformations to apply upon
↳ loading.
        """
        self.paths = paths
        self.transform = transforms
        self.op = op

        self.num_ex = len(paths)

        try:
            if self.op == 'train':
                self.split_paths = self.paths[0:int(0.6*self.num_ex)]
            elif self.op == 'val':
                self.split_paths = self.paths[int(0.6*self.num_ex):int(0.8*self.
↳ num_ex)]
            elif self.op == 'test':
                self.split_paths = self.paths[int(0.8*self.num_ex):int(self.
↳ num_ex)]

```

```

    except ValueError:
        print('op is not train, val, or test')

def __len__(self):
    return len(self.split_paths)

def __getitem__(self, idx):
    dmap = io.imread(self.split_paths[idx][1])
    img = io.imread(self.split_paths[idx][0])

    img = resize(img, (img.shape[0] // 2, img.shape[1] // 2),
                  anti_aliasing=True)
    dmap = resize(dmap, (dmap.shape[0] // 2, dmap.shape[1] // 2),
                  anti_aliasing=True)
    dmap = np.expand_dims(rgb2gray(dmap), 0)
    if self.transform:
        img, dmap = self.img_transform(img, dmap)

    sample = {'img': img, 'dmap': dmap}
    return sample

def img_transform(self, img, dmap):
    ## Apply Transformations
    img = self.transform(img)
    dmap = torch.from_numpy(dmap).type(torch.float32)
    return img, dmap

```

```

[ ]: data_paths = getDataPaths()
img_transform = transforms.Compose([
    transforms.ToTensor(),
])

train_dataset = ImageDataset(data_paths, op="train", transforms=img_transform)
val_dataset = ImageDataset(data_paths, op="val", transforms=img_transform)
test_dataset = ImageDataset(data_paths, op="test", transforms=img_transform)

print(train_dataset.__len__())
print(val_dataset.__len__())
print(test_dataset.__len__())

train_batch_size = 1
val_batch_size = 1
test_batch_size = 1

train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size,
    ↪shuffle=False)

```

```

validation_dataloader = DataLoader(val_dataset, batch_size=val_batch_size,
    ↪shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)

train_dataset_notensor = ImageDataset(data_paths, op="train", transforms=None)
train_dataloader_notensor = DataLoader(train_dataset_notensor,
    ↪batch_size=train_batch_size, shuffle=False)

# Load some image and view them
for k,data in enumerate(train_dataloader_notensor):
    plt.figure()
    plt.imshow(data['img'][0])
    plt.figure()
    plt.imshow(data['dmap'][0][0])
    plt.show()
    if k >= 3:
        break

```

```

[ ]: ### UNET FROM https://pytorch.org/hub/
    ↪mateuszbuda_brain-segmentation-pytorch_unet/
from collections import OrderedDict
import torch
import torch.nn as nn

### Model definitions.

class pretrainedG():
    def __init__(self,in_channels, out_channels, init_features):
        self.in_channels = in_channels;
        self.out_channels = out_channels;
        self.init_features = init_features;
    def get(self):
        gM = torch.hub.load('mateuszbuda/brain-segmentation-pytorch', 'UNET',
            in_channels=self.in_channels, out_channels=self.out_channels,
            init_features=self.init_features, pretrained=True)
        return gM

class Discriminator(nn.Module): # The PatchGAN discriminator
    def __init__(self, in_channels=1):
        super(Discriminator, self).__init__()

        def discriminator_block(in_filters, out_filters, normalization=True):
            """Returns downsampling layers of each discriminator block"""
            layers = [nn.Conv2d(in_filters, out_filters, 4, stride=2,
    ↪padding=1)]
            if normalization:
                layers.append(nn.InstanceNorm2d(out_filters))

```

```

        layers.append(nn.LeakyReLU(0.2, inplace=True))
    return layers

    self.gM = nn.Sequential(
        *discriminator_block(in_channels * 2, 64, normalization=False),
        *discriminator_block(64, 128),
        *discriminator_block(128, 256),
        *discriminator_block(256, 512),
        nn.ZeroPad2d((1, 0, 1, 0)),
        nn.Conv2d(512, 1, 4, padding=1, bias=False)
    )

    def forward(self, img_A, img_B):
        # Concatenate image and condition image by channels to produce input
        img_input = torch.cat((img_A, img_B), 1)
        return self.gM(img_input)

class pretrainedD(): # Pretrained ResNet 18 as discriminator (Not Used)
    def __init__(self, in_channels, out_channels, init_features):
        self.in_channels = in_channels;
        self.out_channels = out_channels;
        self.init_features = init_features;
    def get(self):
        model = torch.hub.load('pytorch/vision', 'resnet18', pretrained=True)
        return model
patch = (1, 128 //16, 416 //16)

```

```
[ ]: print(*patch)
```

```

[ ]: ## Initialize Optimizer and Learning Rate Scheduler

EPOCHS = 30
VISUALIZE = True
TRAIN = 1
MODEL_SAVE_DIR = r'/media/athrva/New Volume/cis520/Final Project/TrainedGAN/'
MODEL_SAVE_DIR_RUD = r'/media/athrva/New Volume/cis520/Final Project/GAN_MODELS/'
↳
TRAIN_LOG_DIR_RUD = r'/media/athrva/New Volume/cis520/Final Project/'
TRAIN_LOG_DIR = r'/media/athrva/New Volume/cis520/Final Project/'

learning_rate = 1e-4
num_epochs = 1

loss_metric = torch.nn.MSELoss() # For regression task

if torch.cuda.is_available():

```

```

    gpu_boole = True
else:
    gpu_boole = False

```

```

[ ]: # get logger

use_cuda = torch.cuda.is_available()

trainLogger = open(TRAIN_LOG_DIR + 'train.log', 'a+') # Logger

def sc1(tensor):
    tensor = (tensor[0,0,:,:] + tensor[0,1,:,:] + tensor[0,2,:,:])/3.

    return tensor.unsqueeze(0).unsqueeze(0)

cuda = True if torch.cuda.is_available() else False

Tensor = torch.cuda.FloatTensor if cuda else torch.FloatTensor # Generic
↳ definition of Tensor default type

if TRAIN == 1:
    print('Pre-trained Model Found...Reusing')

    gM = pretrainedG(in_channels=3, out_channels= 1, init_features= 32)
    gM = gM.get()
    dM = Discriminator(in_channels = 1)
    if use_cuda:
        gM.cuda()
        dM.cuda()
    device = torch.device('cuda')
    if os.path.exists(MODEL_SAVE_DIR + '/SegNet.pt'):
        checkpoint = torch.load(MODEL_SAVE_DIR + '/SegNet.pt',
↳map_location=device)
        gM.load_state_dict(checkpoint['model_state_dict'])
        dM.load_state_dict(checkpoint['dis_state_dict'])
        optimizer_G = torch.optim.Adam(gM.parameters(), lr=learning_rate)
        optimizer_D = torch.optim.Adam(dM.parameters(), lr=learning_rate/2) #
↳Half learning rate for the discriminator
        optimizer_G.load_state_dict(checkpoint['G_optimizer_state_dict'])
        optimizer_D.load_state_dict(checkpoint['D_optimizer_state_dict'])
        schedulerG = torch.optim.lr_scheduler.StepLR(optimizer_G, step_size=10,
↳gamma=0.1)
        schedulerD = torch.optim.lr_scheduler.StepLR(optimizer_D, step_size=10,
↳gamma=0.1)

```



```

        dM.train()
        gM.train()
    else:
        if use_cuda:
            gM.cuda()
            dM.cuda()

            optimizer_G = torch.optim.Adam(gM.parameters(), lr=learning_rate)
            optimizer_D = torch.optim.Adam(dM.parameters(), lr=learning_rate/2)
            schedulerG = torch.optim.lr_scheduler.StepLR(optimizer_G, step_size=10,
↪gamma=0.1)
            schedulerD = torch.optim.lr_scheduler.StepLR(optimizer_D, step_size=10,
↪gamma=0.1)

        dM.train()
        gM.train()

    # similarity loss definition
    loss_fn = torch.nn.CrossEntropyLoss()
    criterion_GAN = torch.nn.L1Loss()

    # continuity loss definition
    loss_hpy = torch.nn.L1Loss(size_average = True)
    loss_hpy_ = torch.nn.MSELoss()

    a = torch.Tensor([[1, 0, -1],
        [2, 0, -2],
        [1, 0, -1]]) # Kernel for gradient computation

    a = a.view((1,1,3,3)).cuda()

    b = torch.Tensor([[1, 2, 1],
        [0, 0, 0],
        [-1, -2, -1]]) # Kernel for gradient computation

    b = b.view((1,1,3,3)).cuda()

    sig = torch.nn.Sigmoid()

    for epoch in range(EPOCHS):
        for batch_idx, batch in enumerate(train_dataloader):

            x = batch["img"]
            y = batch["dmap"]
            if gpu_boole:
                x = x.cuda().float()

```

```

        y = y.cuda().float()

    epoch_loss = 0

    real_A, real_B = x, y
    real_B = real_B.cuda()
    real_A = real_A.cuda()
    real_B = real_B

    optimizer_G.zero_grad()

    fake_B = gM(real_A)
    fake_B = sig(fake_B) # pass the model output through a sigmoid
    ↪ (optional)

    # Adversarial ground truths
    valid = Variable(Tensor(np.ones((real_A.size(0), *patch))),
    ↪ requires_grad=False)
    fake = Variable(Tensor(np.zeros((real_A.size(0), *patch))),
    ↪ requires_grad=False)

    pred_fake = dM(fake_B, sc1(real_A))

    """Code for the novel gradient loss"""

    G_x = F.conv2d(real_B, a)
    G_y = F.conv2d(real_B, b)
    F_x = F.conv2d(fake_B, a)
    F_y = F.conv2d(fake_B, b)
    G = sig((torch.pow(G_x,2)/torch.max(torch.pow(G_x,2)))+(torch.
    ↪ pow(G_y,2)/torch.max(torch.pow(G_y,2))))
    F_ = sig((torch.pow(F_x,2)/torch.max(torch.pow(F_x,2)))+(torch.
    ↪ pow(F_y,2)/torch.max(torch.pow(F_y,2))))

    lhpy = loss_hpy(fake_B,real_B)
    lhpy_ = loss_hpy_(fake_B,real_B)

    # Total loss : It has a few components
    loss_G = 0.5*torch.sum(torch.abs(torch.abs(G) - torch.abs(F_))) +
    ↪ 1*(lhpy + lhpy_) + 0.5*loss_fn(pred_fake, valid) +
    ↪ 3*criterion_GAN(pred_fake, valid)

    loss_G.backward()

```

```

epoch_loss += loss_G.item()
optimizer_G.step()

    if batch_idx % 5 == 0: # Update the discriminator every 5 batches
↳ instead of 1.
        optimizer_D.zero_grad()
        #pass
        # Real loss
        pred_real = dM(real_B, sc1(real_A))
        loss_real = criterion_GAN(pred_real, valid)

        # Fake loss
        pred_fake = dM(fake_B.detach(), sc1(real_A))
        loss_fake = criterion_GAN(pred_fake, fake)

        # Total loss
        loss_D = (loss_real + loss_fake)
        loss_D.backward()
        optimizer_D.step()

trainLogger.write(f'{loss_G.item()}'+', '+f'{loss_D.item()}'+'\n')
with open(TRAIN_LOG_DIR_RUD + 'training_info.txt', 'a+') as file: #
↳ redundant logger
    file.write(f'G Loss : {loss_G.item()}, D Loss : {loss_D.
↳ item()}')

    if VISUALIZE:
        im_target = fake_B.detach().cpu().numpy()[0]
        im_target = im_target/abs(im_target.max())
        im_targetgt = real_B.detach().cpu().numpy()[0]
        im_target_rgb = im_target
        im_real = x[0].detach().cpu().numpy()
        im_target_rgb = np.vstack([rgb2gray(np.
↳ moveaxis(im_real, 0, -1)), im_targetgt[0],
                                im_target_rgb[0]])
        cv2.imshow('Output : ', im_target_rgb)
        cv2.waitKey(5)

    print('Current Epoch Loss : ', epoch_loss)

    print(batch_idx, '/', EPOCHS, '|', ' | loss G :', loss_G.item(), '|',
↳ '| loss D :', loss_D.item())

    torch.save({'model_state_dict' : gM.state_dict(),
                'G_optimizer_state_dict' : optimizer_G.state_dict(),
                'D_optimizer_state_dict' : optimizer_D.state_dict(),

```

```

        'dis_state_dict' : dM.state_dict()}, MODEL_SAVE_DIR_RUD,
↪+ f'SegNet_17{epoch}.pt');
    torch.save({'model_state_dict' : gM.state_dict(),
        'G_optimizer_state_dict' : optimizer_G.state_dict(),
        'D_optimizer_state_dict' : optimizer_D.state_dict(),
        'dis_state_dict' : dM.state_dict()}, MODEL_SAVE_DIR +
↪'SegNet.pt');

```

```

[ ]: ## PREDICTION PIPELINE.
# This is a simple script to visualize the model's performance on the data
↪images.

use_cuda = torch.cuda.is_available()
PREDICT = 1

once = True

def scl(tensor):
    tensor = (tensor[0,0,:,:] + tensor[0,1,:,:] + tensor[0,2,:,:])/3.
    return tensor.unsqueeze(0).unsqueeze(0)

cuda = True if torch.cuda.is_available() else False

Tensor = torch.cuda.FloatTensor if cuda else torch.FloatTensor

if PREDICT == 1:
    print('Pre-trained Model Found...Reusing')

    gM = pretrainedG(in_channels=3, out_channels= 1, init_features= 32)
    gM = gM.get()
    dM = Discriminator(in_channels = 1)
    if use_cuda:
        gM.cuda()
        dM.cuda()
    device = torch.device('cuda')
    if os.path.exists(MODEL_SAVE_DIR + '/SegNet.pt'):
        checkpoint = torch.load(MODEL_SAVE_DIR + '/SegNet.pt',
↪map_location=device)
        gM.load_state_dict(checkpoint['model_state_dict'])
        dM.load_state_dict(checkpoint['dis_state_dict'])
        optimizer_G = torch.optim.Adam(gM.parameters(), lr=learning_rate)
        optimizer_D = torch.optim.Adam(dM.parameters(), lr=learning_rate/2)
        optimizer_G.load_state_dict(checkpoint['G_optimizer_state_dict'])
        optimizer_D.load_state_dict(checkpoint['D_optimizer_state_dict'])
        schedulerG = torch.optim.lr_scheduler.StepLR(optimizer_G, step_size=10,
↪gamma=0.1)

```

```

        schedulerD = torch.optim.lr_scheduler.StepLR(optimizer_D, step_size=10,
↪gamma=0.1)
        dM.train()
        gM.train()
    else:
        print('No Model Found...')
        dM.train()
        gM.train()

sig = torch.nn.Sigmoid()

for epoch in range(1):
    for batch_idx, batch in enumerate(validation_dataloader): # could be
↪test_dataloader.

        x = batch["img"]
        y = batch["dmap"]
        if gpu_boole:
            x = x.cuda().float()
            y = y.cuda().float()

        epoch_loss = 0

        real_A, real_B = x, y
        real_B = real_B.cuda()
        real_A = real_A.cuda()
        real_B = real_B

        optimizer_G.zero_grad()

        fake_B = gM(real_A)
        im_target = fake_B.detach().cpu().numpy()[0]
        im_target = im_target/abs(im_target.max())
        im_targetgt = real_B.detach().cpu().numpy()[0]
        im_target_rgb = im_target
        im_real = x[0].detach().cpu().numpy()
        im_target_rgb = np.vstack([im_targetgt[0],
                                   im_target_rgb[0]])

        if once == True:
            plt.imsave(basePath + os.path.sep + 'output.png', im_target_rgb)
            once = False
        cv2.imshow('Output : ', im_target_rgb)
        cv2.waitKey(5)
        plt.imsave(f'/media/athrva/New Volume/cis520/Final Project/
↪Progression/out_{batch_idx}.png', im_target_rgb)

```

